



MITx 6.86x

Machine Learning with Python-From Linear Models to Deep Learning

[Help](#)

smitha_kannur ▾

[Course](#)

[Progress](#)

[Dates](#)

[Discussion](#)

[Resources](#)

[Home](#) [Course](#) / [Unit 3 Neural networks \(2.5 weeks\)](#) / [Homework 3](#)

[< Previous](#)



[Next >](#)


3. Backpropagation

[Bookmark this page](#)

Homework due Oct 28, 2020 05:29 IST *Completed*

One of the key steps for training multi-layer neural networks is stochastic gradient descent. We will use the back-propagation algorithm to compute the gradient of the loss function with respect to the model parameters.

Consider the L -layer neural network below:



In the following problems, we will use the following notation: b_j^l is the bias of the j^{th} neuron in the l^{th} layer, a_j^l is the activation of j^{th} neuron in the l^{th} layer, and w_{jk}^l is the weight for the connection from the k^{th} neuron in the $(l-1)^{th}$ layer to the j^{th} neuron in the l^{th} layer.

If the activation function is f and the loss function we are minimizing is C , then the equations describing the network are:

$$a_j^l = f\left(\sum_k w_{jk}^l a_k^{l-1} + b_j^l\right)$$

$$\text{Loss} = C(a^L)$$

Note that notations without subscript denote the corresponding vector or matrix, so that a^l is activation vector of the l^{th} layer, and w^l is the weights matrix in l^{th} layer.

For $l = 1, \dots, L$.

Computing the Error

2/2 points (graded)

Let the weighted inputs to the d neurons in layer l be defined as $z^l \equiv w^l a^{l-1} + b^l$, where $z^l \in \mathbb{R}^d$. As a result, we can also write the activation of layer l as $a^l \equiv f(z^l)$, and the "error" of neuron j in layer l as $\delta_j^l \equiv \frac{\partial C}{\partial z_j^l}$. Let $\delta^l \in \mathbb{R}^d$ denote the full vector of errors associated with layer l .

Back-propagation will give us a way of computing δ^l for every layer.

Assume there are d outputs from the last layer (i.e. $a^L \in \mathbb{R}^d$). What is δ_j^L for the last layer?

☒ $\frac{\partial C}{\partial a_j^L} f'(z_j^L)$

☐ $\sum_{k=1}^d \frac{\partial C}{\partial a_k^L} f'(z_j^L)$

☐ $\frac{\partial C}{\partial a_j^L}$

☐ $f'(z_j^L)$



What is δ_j^l for all $l \neq L$?

☒ $\sum_k w_{kj}^{l+1} \delta_k^{l+1} f'(z_j^l)$

☐ $\delta_k^{l+1} f'(z_j^l)$

☐ $\sum_k w_{jk}^{l-1} \delta_j^{l-1} f'(z_j^l)$

☐ $\sum_k w_{kj}^{l+1} \delta_k^{l+1} f(z_j^l)$



Solution:

We make use of the chain rule.

1. By definition, $\delta_j^L = \frac{\partial C}{\partial a_j^L} \frac{\partial a_j^L}{\partial z_j^L} = \frac{\partial C}{\partial a_j^L} f'(z_j^L)$.

2. We have:

$$\begin{aligned} \delta_j^l &= \frac{\partial C}{\partial z_j^l} \\ &= \sum_k \frac{\partial C}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} \\ &= \sum_k \frac{\partial z_k^{l+1}}{\partial z_j^l} \delta_k^{l+1} \end{aligned}$$

Then we have $z_k^{l+1} = \sum_j w_{kj}^{l+1} a_j^l + b_k^{l+1} = \sum_j w_{kj}^{l+1} f(z_j^l) + b_k^{l+1}$. Taking the derivative of this with respect to z_j^l gives $w_{kj}^{l+1} f'(z_j^l)$.

Combining the two gives the final answer: $\delta_j^l = \sum_k w_{kj}^{l+1} \delta_k^{l+1} f'(z_j^l)$.

Submit

You have used 2 of 2 attempts

i Answers are displayed within the problem

Parameter Derivatives

2/2 points (graded)

During SGD we are interested in relating the errors computed by back-propagation to the quantities of real interest: the partial derivatives of the loss with respect to our parameters. Here that is $\frac{\partial C}{\partial w_{jk}^l}$ and $\frac{\partial C}{\partial b_j^l}$.

What is $\frac{\partial C}{\partial w_{jk}^l}$? Write in terms of the variables a_k^{l-1} , w_j^l , b_j^l , and δ_j^l if necessary.

Example of writing superscripts and subscripts:

δ_j^{l+1} for δ_j^l

w_{jk}^{l-1} for w_{jk}^l

$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \cdot \delta_j^l$$

✓ Answer: $a_k^{l-1} \cdot \delta_j^l$

What is $\frac{\partial C}{\partial b_j^l}$? Write in terms of the variables a_k^{l-1} , w_{jk}^l , b_j^l , and δ_j^l if necessary.

$$\frac{\partial C}{\partial b_j^l} = \delta_j^l$$

✓ Answer: δ_j^l

STANDARD NOTATION

Solution:

$$1. \frac{\partial C}{\partial w_{jk}^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$$

$$2. \frac{\partial C}{\partial b_j^l} = \frac{\partial C}{\partial z_j^l} \frac{\partial z_j^l}{\partial b_j^l} = 1 \cdot \delta_j^l$$

Submit

You have used 1 of 5 attempts

🔔 Answers are displayed within the problem

Activation Functions: Sigmoid

4/4 points (graded)

Recall that there are several different possible choices of activation functions f . Let's get more familiar with them and their gradients.

What is the derivative of the sigmoid function, $\sigma(z) = \frac{1}{1+e^{-z}}$? Please write your answer in terms of e and z :

$$(1 - \frac{1}{1+e^{-z}}) \cdot \frac{1}{1+e^{-z}}$$

✓ Answer: $e^{-z} / (1 + e^{-z})^2$

$$\left(1 - \frac{1}{1+e^{-z}}\right) \cdot \left(\frac{1}{1+e^{-z}}\right)$$

Which of the following is true of $\sigma'(z)$ as $||z||$ gets large?

☐ Its magnitude becomes large.

☒ Its magnitude becomes small.

☐ It suffers from high variance.



What is the derivative of the ReLU function, $\text{ReLU}(z) = \max(0, z)$ for $z > 0$?

1

✓ Answer: 1

1

For $z < 0$?

✓ Answer: 0

[STANDARD NOTATION](#)**Solution:**

$\sigma'(z) = \sigma(z)(1 - \sigma(z))$. As z gets large in magnitude, the sigmoid function saturates, and the gradient approaches zero.

ReLU is a simple activation function. Above zero, it has a constant gradient of 1. Below zero, it is always zero.

You have used 1 of 5 attempts

i Answers are displayed within the problem

Simple Network

0/4 points (graded)

Consider a simple 2-layer neural network with a single neuron in each layer. The loss function is the quadratic loss:

$$C = \frac{1}{2}(y - t)^2, \text{ where } y \text{ is the prediction and } t \text{ is the target.}$$

Starting with input x we have:

- $z_1 = w_1 x$
- $a_1 = \text{ReLU}(z_1)$
- $z_2 = w_2 a_1 + b$
- $y = \sigma(z_2)$
- $C = \frac{1}{2}(y - t)^2$

Consider a target value $t = 1$ and input value $x = 3$. The weights and bias are $w_1 = 0.01$, $w_2 = -5$, and $b = -1$.

Please provide numerical answers accurate to at least three decimal places.

What is the loss?

✗ Answer: 0.28842841648243966

What are the derivatives with respect to the parameters?

$$\frac{\partial C}{\partial w_1} =$$

✗ Answer: 2.0809165621704553

$$\frac{\partial C}{\partial w_2} =$$

✗ Answer: -0.00416183312434091

$$\frac{\partial C}{\partial b} =$$

✗ Answer: -0.13872777081136367

[STANDARD NOTATION](#)**Solution:**

Using the chain rule, we have:

- $\frac{\partial C}{\partial w_1} = \frac{\partial C}{\partial y} \frac{\partial y}{\partial z_2} \frac{\partial z_2}{\partial a_1} \frac{\partial a_1}{\partial z_1} \frac{\partial z_1}{\partial w_1} = (y - t) y (1 - y) w_2 \mathbf{1}\{z_1 > 0\} x$
- $\frac{\partial C}{\partial w_2} = \frac{\partial C}{\partial y} \frac{\partial y}{\partial z_2} \frac{\partial z_2}{\partial w_2} = (y - t) y (1 - y) a_1$
- $\frac{\partial C}{\partial b} = (y - t) y (1 - y)$

Submit

You have used 1 of 5 attempts

i Answers are displayed within the problem

SGD

1/1 point (graded)

Referring to the previous problem, what is the update rule for w_1 in the SGD algorithm with step size η ? Write in terms of w_1 , η , and $\frac{\partial C}{\partial w_1}$; enter the latter as (partialC)/(partialw_1), noting the lack of space in the variable names:

Next $w_1 =$ **✓ Answer:** `w_1 - eta * (partialC)/(partialw_1)`

STANDARD NOTATION

Solution:

The definition of the simple SGD update rule is `new_parameter = old_parameter - learning_rate * derivative of loss w.r.t old parameter`.

Submit

You have used 1 of 5 attempts

i Answers are displayed within the problem

Discussion

Hide Discussion

Topic: Unit 3 Neural networks (2.5 weeks):Homework 3 / 3. Backpropagation

Add a Post

Show all posts ▼

by recent activity ▼

- [Another approach - Simple Network](#) 12
 Pinned Community TA
- [Backpropagation resource - Simple network](#) 7
I found this well done and helpful to apply backpropagation, also in the context of the assignment "Simple Network": <https://www.y...>
 Pinned
- [On the notation](#) 4
For clarity, I believe the weight sub-indexes could be switched to make it consistent with regular matrix notation. For example, $W_{Lj, \dots}$
- [What is the point of doing an online course in deep learning when you have to derive and understand the equations on your own?](#) 25
I am probably missing the point of the course structure but I don't understand what is the point of basically assigning to you the tas...
- [Staff:about 'Sorry, couldn't parse formula'](#) 7
even if i am entering the formula as per given technique, it says 'Sorry, couldn't parse formula'.
- [Simple Network chain derivatives answer input](#) 3
Hi everyone, Once again I find myself loosing points due to not understanding the right way to input the solution (happened already...

💬 [STAFF] Simple Network Help	2
[STAFF] Im stuck on this question. I tried to start by taking the derivatives for y, but y is defined in terms of sigmoid, wich is defined i...	
💬 Hint for Parameter Derivatives	4
I found it a lot easier to do the Simple Network question first, and then I went back and did Parameter Derivatives. I think I could hav...	
? Simple Network	17
I am not able to get the derivative of the loss with respect to w 1 even though i could get all the other answers correct. Not sure whe...	
? Simple Network: does it make sense to expand?	6
✓ SGD: how to represent η in standard notation?	3
SGD: how to represent η in standard notation?	
💬 On the 'simple network' section	1
One way and less confusing for me is just to use the chain rule on differentating the individual expressions given for z1, a1, z2, y and...	
💬 [STAFF] "Parameter Derivatives": question formatting issue delta i\A	4

< Previous

Next >

© All Rights Reserved



edX

[About](#)

[Affiliates](#)

[edX for Business](#)

[Open edX](#)

[Careers](#)

[News](#)

Legal

[Terms of Service & Honor Code](#)

[Privacy Policy](#)

[Accessibility Policy](#)
[Trademark Policy](#)
[Sitemap](#)

Connect

[Blog](#)
[Contact Us](#)
[Help Center](#)
[Media Kit](#)
[Donate](#)



© 2020 edX Inc. All rights reserved.

深圳市恒宇博科技有限公司 [粤ICP备17044299号-2](#)