



MITx 6.86x

[Help](#)

smitha\_kannur ▾

[Course](#)

[Progress](#)

[Dates](#)

[Discussion](#)

[Resources](#)

[Home](#) / [Course](#) / [Unit 2 Nonlinear Classification, Linear regression, Collaborative Filtering \(2 weeks\)](#) / [Project 2: Digit recognition \(Part 1\)](#)

[< Previous](#)



[Next >](#)

## 3. Support Vector Machine

[Bookmark this page](#)

Project due Oct 21, 2020 05:29 IST *Completed*

Bob thinks it is clearly not a regression problem, but a classification problem. He thinks that we can change it into a binary classification and use the support vector machine we learned in Lecture 4 to solve the problem. In order to do so, he suggests that we can build an one vs. rest model for every digit. For example, classifying the digits into two classes: 0 and not 0.

Bob wrote a function `run_svm_one_vs_rest_on_MNIST` where he changed the labels of digits 1-9 to 1 and keeps the label 0 for digit 0. He also found that `sklearn` package contains an SVM model that you can use directly. He gave you the link to this model and hopes you can tell him how to use that.

**You will be working in the file `part1/svm.py` in this problem**

**Important:** For this problem, you will need to use the `scikit-learn` library. If you don't have it, install it using `pip install sklearn`

## One vs. Rest SVM

5.0/5.0 points (graded)

Use the `sklearn` package and build the SVM model on your local machine. Use `random_state = 0`, `C=0.1` and default values for other parameters.

**Available Functions:** You have access to the `sklearn`'s implementation of the linear SVM as `LinearSVC`; No need to import anything.

```
1 def one_vs_rest_svm(train_x, train_y, test_x):
2     """
3     Trains a linear SVM for binary classification
4
5     Args:
6         train_x - (n, d) NumPy array (n datapoints each with d features)
7         train_y - (n, ) NumPy array containing the labels (0 or 1) for each training data point
8         test_x - (m, d) NumPy array (m datapoints each with d features)
9     Returns:
10        pred_test_y - (m,) NumPy array containing the labels (0 or 1) for each test data point
11    """
12    try:
13        svm = LinearSVC(random_state=0, C=0.1)
14        svm.fit(train_x, train_y)
15        pred_test_y = svm.predict(test_x)
16    return pred_test_y
```

Press ESC then TAB or click outside of the code editor to exit

Correct

## Test results

CORRECT

[See full output](#)

[See full output](#)

Submit

You have used 2 of 25 attempts

## Binary classification error

5.0/5.0 points (graded)

Report the test error by running `run_svm_one_vs_rest_on_MNIST`.

Error = 0.007499999999999951 ✓ Answer: 0.007499999999999951

Submit

You have used 1 of 20 attempts

Answers are displayed within the problem

## Implement C-SVM

0.0/5.0 points (graded)

Play with the C parameter of SVM, what statement is true about the C parameter?

(Choose all that apply.)

☒ Larger C gives larger tolerance of violation.

☐ Larger C gives smaller tolerance of violation. ✓

☐ Larger C gives a larger-margin separating hyperplane.

☒ Larger C gives a smaller-margin separating hyperplane. ✓

✗

### Solution:

C represents the tolerance of error. A larger C means we are punishing more on the classification error, thus being less tolerant to misclassifications. Therefore, we will get a smaller margin hyperplane.

Submit

You have used 2 of 2 attempts

Answers are displayed within the problem

## Multiclass SVM

5.0/5.0 points (graded)

In fact, `sklearn` already implements a multiclass SVM with a one-vs-rest strategy. Use `LinearSVC` to build a multiclass SVM model

**Available Functions:** You have access to the sklearn's implementation of the linear SVM as `LinearSVC`; No need to import anything.

```
1 def multi_class_svm(train_x, train_y, test_x):
2     """
3     Trains a linear SVM for multiclass classification using a one-vs-rest strategy
4
5     Args:
6         train_x - (n, d) NumPy array (n datapoints each with d features)
7         train_y - (n, ) NumPy array containing the labels (int) for each training data point
8         test_x - (m, d) NumPy array (m datapoints each with d features)
9     Returns:
10        pred_test_y - (m,) NumPy array containing the labels (int) for each test data point
11    """
12    try:
13        svm = LinearSVC(random_state=0, C=0.1)
14        svm.fit(train_x, train_y)
15        pred_test_y = svm.predict(test_x)
16    return pred_test_y
```

Press ESC then TAB or click outside of the code editor to exit

Correct



<p>How to interpret the test result?</p> <p>How to interpret the test result? why following are correct? ## one vs rest svm case train_y: [0 0 1 0 0 1 0 0 1 0 0 0 1 1 1 1] Submis...</p>	1
<p>Tolerance of Violation</p> <p>Is this a technical term with a specific meaning?</p>	13
<p>Nice article on "The Math behind Linear SVC Classifier"</p> <p>https://www.kaggle.com/yinguawang/the-math-behind-linear-svc-classifier</p>	1

< Previous

Next >

© All Rights Reserved



## edX

[About](#)  
[Affiliates](#)  
[edX for Business](#)  
[Open edX](#)  
[Careers](#)  
[News](#)

## Legal

[Terms of Service & Honor Code](#)  
[Privacy Policy](#)  
[Accessibility Policy](#)  
[Trademark Policy](#)  
[Sitemap](#)

## Connect

[Blog](#)  
[Contact Us](#)  
[Help Center](#)  
[Media Kit](#)  
[Donate](#)



© 2020 edX Inc. All rights reserved.

深圳市恒宇博科技有限公司 [粤ICP备17044299号-2](#)